

Bank Churn Prediction

By Raghuram Palaniappan

Problem Statement

Context

Businesses like banks which provide service have to worry about problem of 'Customer Churn' i.e. customers leaving and joining another service provider. It is important to understand which aspects of the service influence a customer's decision in this regard. Management can concentrate efforts on improvement of service, keeping in mind these priorities.

Objective

You as a Data scientist with the bank need to build a neural network based classifier that can determine whether a customer will leave the bank or not in the next 6 months.

Data Dictionary

- CustomerId: Unique ID which is assigned to each customer
- Surname: Last name of the customer
- CreditScore: It defines the credit history of the customer.
- Geography: A customer's location
- Gender: It defines the Gender of the customer
- Age: Age of the customer
- Tenure: Number of years for which the customer has been with the bank
- NumOfProducts: refers to the number of products that a customer has purchased through the bank.
- Balance: Account balance
- HasCrCard: It is a categorical variable which decides whether the customer has credit card or not.
- EstimatedSalary: Estimated salary

- **IsActiveMember:** Is is a categorical variable which decides whether the customer is active member of the bank or not (Active member in the sense, using bank products regularly, making transactions etc)
- **Exited :** whether or not the customer left the bank within six month. It can take two values
 - 0 = No (Customer did not leave the bank)
 - 1 = Yes (Customer left the bank)

Importing necessary libraries

```
In [ ]: # Libraries to help with reading and manipulating data
import pandas as pd
import numpy as np

# Libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Library to split data
from sklearn.model_selection import train_test_split

# Library to import to standardize the data
from sklearn.preprocessing import StandardScaler, LabelEncoder

# importing different functions to build models
import tensorflow as tf
from tensorflow import keras
from keras import backend
from keras.models import Sequential
from keras.layers import Dense, Dropout

# importing SMOTE
from imblearn.over_sampling import SMOTE

# importing metrics
from sklearn.metrics import confusion_matrix, roc_curve, classification_report, recall_score

import random

# Library to avoid the warnings
import warnings
warnings.filterwarnings("ignore")
```

WARNING:tensorflow:From C:\Users\Raghuram\AppData\Roaming\Python\Python311\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Loading the dataset

```
In [ ]: ds = pd.read_csv("Churn.csv")
```

```
In [ ]: df = ds.copy()
```

Data Overview

```
In [ ]: ds.head()
```

```
Out[ ]:   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  Tenure  Balance  Nu
```

0	1	15634602	Hargrave	619	France	Female	42	2	0.00
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86
2	3	15619304	Onio	502	France	Female	42	8	159660.80
3	4	15701354	Boni	699	France	Female	39	1	0.00
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82

```
In [ ]: ds.shape
```

```
Out[ ]: (10000, 14)
```

```
In [ ]: ds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [ ]: ds.describe().T
```

Out[]:

	count	mean	std	min	25%	50%	
RowNumber	10000.0	5.000500e+03	2886.895680	1.00	2500.75	5.000500e+03	7.5002!
CustomerId	10000.0	1.569094e+07	71936.186123	15565701.00	15628528.25	1.569074e+07	1.5753!
CreditScore	10000.0	6.505288e+02	96.653299	350.00	584.00	6.520000e+02	7.1800!
Age	10000.0	3.892180e+01	10.487806	18.00	32.00	3.700000e+01	4.4000!
Tenure	10000.0	5.012800e+00	2.892174	0.00	3.00	5.000000e+00	7.0000!
Balance	10000.0	7.648589e+04	62397.405202	0.00	0.00	9.719854e+04	1.2764!
NumOfProducts	10000.0	1.530200e+00	0.581654	1.00	1.00	1.000000e+00	2.0000!
HasCrCard	10000.0	7.055000e-01	0.455840	0.00	0.00	1.000000e+00	1.0000!
IsActiveMember	10000.0	5.151000e-01	0.499797	0.00	0.00	1.000000e+00	1.0000!
EstimatedSalary	10000.0	1.000902e+05	57510.492818	11.58	51002.11	1.001939e+05	1.4938!
Exited	10000.0	2.037000e-01	0.402769	0.00	0.00	0.000000e+00	0.0000!

Checking for Missing Values

In []: `ds.isnull().sum()`

Out[]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

Checking for unique values for each of the column

In []: `ds.nunique()`

```
Out[ ]: RowNumber      10000
        CustomerId    10000
        Surname       2932
        CreditScore    460
        Geography      3
        Gender         2
        Age            70
        Tenure         11
        Balance        6382
        NumOfProducts  4
        HasCrCard      2
        IsActiveMember 2
        EstimatedSalary 9999
        Exited         2
        dtype: int64
```

Observations: The data set is 14 variables and 10,000 rows. There are no missing values. The target dependent variable is the Exited column that determines if a customer left the bank or not within 6 months.

```
In [ ]: #RowNumber , CustomerId and Surname are unique hence dropping it
        ds = ds.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
```

Exploratory Data Analysis

Univariate Analysis

```
In [ ]: # function to plot a boxplot and a histogram along the same scale.

def histogram_boxplot(data, feature, figsize=(12, 7), kde=False, bins=None):
    """
    Boxplot and histogram combined

    data: dataframe
    feature: dataframe column
    figsize: size of figure (default (12,7))
    kde: whether to show the density curve (default False)
    bins: number of bins for histogram (default None)
    """
    f2, (ax_box2, ax_hist2) = plt.subplots(
        nrows=2, # Number of rows of the subplot grid= 2
        sharex=True, # x-axis will be shared among all subplots
        gridspec_kw={"height_ratios": (0.25, 0.75)},
        figsize=figsize,
    ) # creating the 2 subplots
    sns.boxplot(
        data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
    ) # boxplot will be created and a star will indicate the mean value of the column
    sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins, palette="winter"
    ) if bins else sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2
    ) # For histogram
    ax_hist2.axvline(
        data[feature].mean(), color="green", linestyle="--"
```

```

) # Add mean to the histogram
ax_hist2.axvline(
    data[feature].median(), color="black", linestyle="--"
) # Add median to the histogram

```

```

In [ ]: # function to create labeled barplots

def labeled_barplot(data, feature, perc=False, n=None):
    """
    Barplot with percentage at the top

    data: dataframe
    feature: dataframe column
    perc: whether to display percentages instead of count (default is False)
    n: displays the top n category levels (default is None, i.e., display all levels)
    """

    total = len(data[feature]) # length of the column
    count = data[feature].nunique()
    if n is None:
        plt.figure(figsize=(count + 1, 5))
    else:
        plt.figure(figsize=(n + 1, 5))

    plt.xticks(rotation=90, fontsize=15)
    ax = sns.countplot(
        data=data,
        x=feature,
        palette="Paired",
        order=data[feature].value_counts().index[:n].sort_values(),
    )

    for p in ax.patches:
        if perc == True:
            label = "{:.1f}%".format(
                100 * p.get_height() / total
            ) # percentage of each class of the category
        else:
            label = p.get_height() # count of each level of the category

        x = p.get_x() + p.get_width() / 2 # width of the plot
        y = p.get_height() # height of the plot


        ax.annotate(
            label,
            (x, y),
            ha="center",
            va="center",
            size=12,
            xytext=(0, 5),
            textcoords="offset points",
        ) # annotate the percentage

    plt.show() # show the plot

```


Observations on CreditScore: The distribution is close to symmetrical with outliers on the left side of the data.

```
In [ ]: histogram_boxplot(ds, 'CreditScore')
```




Observations on Age: The data is right skewed due with a lot of outliers in the right side of the data.

```
In [ ]: histogram_boxplot(ds, 'Age')
```




Observations on Balance: Outside of the data points around \$0 the distribution is symmetrical. There seem to be a good amount of people with little to no balance in their accounts. The overall data is right skewed.

```
In [ ]: histogram_boxplot(ds, 'Balance')
```




Observations on Estimated Salary: The distribution is pretty uniform. The data is spread relatively even.

```
In [ ]: histogram_boxplot(ds, 'EstimatedSalary')
```




Observations on Exited: Around 20% of customers of the bank have left within the first 6 months.


```
In [ ]: labeled_barplot(ds, "Exited", perc=True)
```




Observations on Geography: Half of the customers come from France with Germany and Spain taking up a quarter each.

```
In [ ]: labeled_barplot(ds, "Geography", perc=True)
```




Observations on Gender: More men at 54.6% are customers of the bank.

```
In [ ]: labeled_barplot(ds, "Gender", perc=True)
```




Observations on Tenure: For 1-9 years, at every year there is around 10% of customers who have that long of a tenure.

```
In [ ]: labeled_barplot(ds, "Tenure", perc=True)
```




Observations on Number of Products: Half of the customers have purchased 1 product through the bank and another 45% have purchased exactly 2 products.

```
In [ ]: labeled_barplot(ds, "NumOfProducts", perc=True)
```




Observations on Has Credit Card: Majority of the customers at 70.5% have credit cards issued by the bank.

```
In [ ]: labeled_barplot(ds, "HasCrCard", perc=True)
```



Observations on Is Active Member: Slightly over half at 51.5% are active customers at the banks.

```
In [ ]: labeled_barplot(ds, "IsActiveMember", perc=True)
```



Bivariate Analysis

```
In [ ]: # function to plot stacked bar chart


def stacked_barplot(data, predictor, target):
    """
    Print the category counts and plot a stacked bar chart

    data: dataframe
    predictor: independent variable
    target: target variable
    """
    count = data[predictor].nunique()
    sorter = data[target].value_counts().index[-1]
    tab1 = pd.crosstab(data[predictor], data[target], margins=True).sort_values(
        by=sorter, ascending=False
    )
    print(tab1)
    print("-" * 120)
    tab = pd.crosstab(data[predictor], data[target], normalize="index").sort_values(
        by=sorter, ascending=False
    )
    tab.plot(kind="bar", stacked=True, figsize=(count + 1, 5))
    plt.legend(
        loc="lower left",
        frameon=False,
    )
    plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
    plt.show()
```

Correlation plot

```
In [ ]: cols_list = ["CreditScore", "Age", "Tenure", "Balance", "EstimatedSalary"]
```

```
In [ ]: plt.figure(figsize=(15, 7))
sns.heatmap(ds[cols_list].corr(), annot=True, vmin=-1, vmax=1, fmt=".2f", cmap="Spectral")
plt.show()
```




Observation: These variables are excellent due to practically no correlation between each other. This shows no multicollinearity.

Exited Vs Geography

```
In [ ]: stacked_barplot(ds, "Geography", "Exited")
```


Exited	0	1	All
Geography			
All	7963	2037	10000
Germany	1695	814	2509
France	4204	810	5014
Spain	2064	413	2477



Exited Vs Gender

```
In [ ]: stacked_barplot(ds, "Gender", "Exited")
```


Exited	0	1	All
Gender			
All	7963	2037	10000
Female	3404	1139	4543
Male	4559	898	5457



Exited Vs Has Credit Card

```
In [ ]: stacked_barplot(ds, "HasCrCard", "Exited")
```


Exited	0	1	All
HasCrCard			
All	7963	2037	10000
1	5631	1424	7055
0	2332	613	2945



Exited Vs Active Member


```
In [ ]: stacked_barplot(ds, "IsActiveMember", "Exited")
```

Exited	0	1	All
IsActiveMember			
All	7963	2037	10000
0	3547	1302	4849
1	4416	735	5151




Exited Vs Credit Score

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='CreditScore',x='Exited',data=ds)
plt.show()
```




Exited Vs Age

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='Age',x='Exited',data=ds)
plt.show()
```




Exited Vs Tenure

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='Tenure',x='Exited',data=ds)
plt.show()
```




Exited Vs Balance

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='Balance',x='Exited',data=ds)
plt.show()
```




Exited Vs Number of Products

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='NumOfProducts',x='Exited',data=ds)
plt.show()
```



Exited Vs Estimated Salary

```
In [ ]: plt.figure(figsize=(5,5))
sns.boxplot(y='EstimatedSalary',x='Exited',data=ds)
plt.show()
```

Observations:

- Geography: The highest ratio of exited customers are from Germany. This variable is a good predictor.
- Gender: The majority of exited customers are females, hence this variable also a good predictor.
- IsActiveMember: A higher ratio of exited customers is observed for non-active customers compared to active customers. This is a good predictor.

EDA Observations

- From the Univariate Analysis we conclude:
 - All bank customers come are based in France, Germany or Spain.
 - The majority of customers are from France, are Male, and have not exited the bank
- From the Bivariate Analysis we conclude:
 - There is no multicollinearity between features.
 - Good predictors: Geography, Gender, and IsActiveMember

Data Preprocessing

Dummy Variable Creation

```
In [ ]: ds = pd.get_dummies(ds, columns=ds.select_dtypes(include=["object"]).columns.tolist(), c
```

Train-validation-test Split

```
In [ ]: X = ds.drop(['Exited'], axis=1)
        y = ds['Exited']
```

```
In [ ]: # Splitting the dataset into the Training and Testing set.

        X_large, X_test, y_large, y_test = train_test_split(X, y, test_size = 0.2, random_stat
```

```
In [ ]: # Splitting the dataset into the Training and Testing set.

        X_train, X_val, y_train, y_val = train_test_split(X_large, y_large, test_size = 0.2, r
```

```
In [ ]: print(X_train.shape, X_val.shape, X_test.shape)

        (6400, 11) (1600, 11) (2000, 11)
```

```
In [ ]: print(y_train.shape, y_val.shape, y_test.shape)

        (6400,) (1600,) (2000,)
```

Data Normalization

Since all the numerical values are on a different scale, so we will be scaling all the numerical values to bring them to the same scale.

```
In [ ]: # creating an instance of the standard scaler
        sc = StandardScaler()

        cols_list = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']

        X_train[cols_list] = sc.fit_transform(X_train[cols_list])
        X_val[cols_list] = sc.transform(X_val[cols_list])
        X_test[cols_list] = sc.transform(X_test[cols_list])
```

Observations: The code has now been split to train test validation sets and the columns have all been scaled properly.

Model Building

Model Evaluation Criterion

Write down the logic for choosing the metric that would be the best metric for this business scenario.

- Our main goal is to reduce false negatives so Recall is our best Metric to select the best models

Let's create a function for plotting the confusion matrix

```
In [ ]: def make_confusion_matrix(actual_targets, predicted_targets):
        """
        To plot the confusion_matrix with percentages

        actual_targets: actual target (dependent) variable values
        predicted_targets: predicted target (dependent) variable values
        """
        cm = confusion_matrix(actual_targets, predicted_targets)
        labels = np.asarray(
            [
                ["{0:0.0f}".format(item) + "\n{0:.2%}".format(item / cm.flatten().sum())]
                for item in cm.flatten()
            ]
        ).reshape(cm.shape[0], cm.shape[1])

        plt.figure(figsize=(6, 4))
        sns.heatmap(cm, annot=labels, fmt="")
        plt.ylabel("True label")
        plt.xlabel("Predicted label")
```

Let's create two blank dataframes that will store the recall values for all the models we build.

```
In [ ]: train_metric_df = pd.DataFrame(columns=["recall"])
        valid_metric_df = pd.DataFrame(columns=["recall"])
```

Neural Network with SGD Optimizer

```
In [ ]: backend.clear_session()
        #Fixing the seed for random number generators so that we can ensure we receive the same
        np.random.seed(2)
        random.seed(2)
        tf.random.set_seed(2)
```

WARNING:tensorflow:From C:\Users\Raghuram\AppData\Roaming\Python\Python311\site-packages\keras\src\backend.py:277: The name tf.reset_default_graph is deprecated. Please use tf.compat.v1.reset_default_graph instead.

```
In [ ]: #Initializing the neural network
        model_0 = Sequential()
        model_0.add(Dense(64, activation='relu', input_dim = X_train.shape[1]))
        model_0.add(Dense(32, activation='relu'))
        model_0.add(Dense(1, activation = 'sigmoid'))
```

```
In [ ]: # SGD Optimizer
        optimizer = tf.keras.optimizers.SGD(0.001)
        metric = keras.metrics.Recall()
```

```
In [ ]: ## Model with binary cross entropy as loss function and recall as the metric.
        model_0.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=[metric])
```

```
In [ ]: model_0.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	768
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
Total params: 2881 (11.25 KB)		
Trainable params: 2881 (11.25 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [ ]: # Fitting the ANN
history_0 = model_0.fit(
    X_train, y_train,
    batch_size=32,
    validation_data=(X_val,y_val),
    epochs=20,
    verbose=1
)
```

Epoch 1/20

WARNING:tensorflow:From C:\Users\Raghuram\AppData\Roaming\Python\Python311\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

200/200 [=====] - 8s 16ms/step - loss: 0.6531 - recall: 0.2124 - val_loss: 0.6103 - val_recall: 0.0429

Epoch 2/20

200/200 [=====] - 2s 9ms/step - loss: 0.5853 - recall: 0.0169 - val_loss: 0.5635 - val_recall: 0.0000e+00

Epoch 3/20

200/200 [=====] - 3s 16ms/step - loss: 0.5473 - recall: 0.0000e+00 - val_loss: 0.5370 - val_recall: 0.0000e+00

Epoch 4/20

200/200 [=====] - 2s 12ms/step - loss: 0.5249 - recall: 0.0000e+00 - val_loss: 0.5210 - val_recall: 0.0000e+00

Epoch 5/20

200/200 [=====] - 1s 4ms/step - loss: 0.5109 - recall: 0.0000e+00 - val_loss: 0.5106 - val_recall: 0.0000e+00

Epoch 6/20

200/200 [=====] - 3s 16ms/step - loss: 0.5015 - recall: 0.0000e+00 - val_loss: 0.5034 - val_recall: 0.0000e+00

Epoch 7/20

200/200 [=====] - 4s 22ms/step - loss: 0.4947 - recall: 0.0000e+00 - val_loss: 0.4980 - val_recall: 0.0000e+00

Epoch 8/20

200/200 [=====] - 4s 22ms/step - loss: 0.4894 - recall: 0.0000e+00 - val_loss: 0.4934 - val_recall: 0.0000e+00

Epoch 9/20

200/200 [=====] - 4s 19ms/step - loss: 0.4849 - recall: 0.0000e+00 - val_loss: 0.4895 - val_recall: 0.0000e+00

Epoch 10/20

200/200 [=====] - 4s 21ms/step - loss: 0.4810 - recall: 0.0000e+00 - val_loss: 0.4860 - val_recall: 0.0000e+00

Epoch 11/20

200/200 [=====] - 3s 16ms/step - loss: 0.4775 - recall: 0.0000e+00 - val_loss: 0.4828 - val_recall: 0.0000e+00

Epoch 12/20

200/200 [=====] - 4s 21ms/step - loss: 0.4742 - recall: 0.0000e+00 - val_loss: 0.4797 - val_recall: 0.0000e+00

Epoch 13/20

200/200 [=====] - 4s 22ms/step - loss: 0.4712 - recall: 0.0000e+00 - val_loss: 0.4769 - val_recall: 0.0000e+00

Epoch 14/20

200/200 [=====] - 4s 22ms/step - loss: 0.4683 - recall: 0.0000e+00 - val_loss: 0.4742 - val_recall: 0.0000e+00

Epoch 15/20

200/200 [=====] - 4s 20ms/step - loss: 0.4656 - recall: 0.0000e+00 - val_loss: 0.4716 - val_recall: 0.0000e+00

Epoch 16/20

200/200 [=====] - 2s 11ms/step - loss: 0.4630 - recall: 0.0000e+00 - val_loss: 0.4691 - val_recall: 0.0000e+00

Epoch 17/20

200/200 [=====] - 3s 15ms/step - loss: 0.4606 - recall: 0.0000e+00 - val_loss: 0.4668 - val_recall: 0.0000e+00

Epoch 18/20

200/200 [=====] - 4s 21ms/step - loss: 0.4582 - recall: 0.0000e+00 - val_loss: 0.4646 - val_recall: 0.0000e+00

Epoch 19/20

200/200 [=====] - 4s 20ms/step - loss: 0.4560 - recall: 0.0000e+00

00e+00 - val_loss: 0.4625 - val_recall: 0.0000e+00


Epoch 20/20

200/200 [=====] - 4s 22ms/step - loss: 0.4539 - recall: 0.00

00e+00 - val_loss: 0.4604 - val_recall: 0.0031


Loss function

```
In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_0.history['loss'])
plt.plot(history_0.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



Recall

```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_0.history['recall'])
plt.plot(history_0.history['val_recall'])
plt.title('model recall')
plt.ylabel('Recall')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
In [ ]: #Predicting the results using best as a threshold
y_train_pred = model_0.predict(X_train)
y_train_pred = (y_train_pred > 0.5)
y_train_pred
```

200/200 [=====] - 4s 13ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [False],
               [False]])
```

```
In [ ]: #Predicting the results using best as a threshold
y_val_pred = model_0.predict(X_val)
y_val_pred = (y_val_pred > 0.5)
y_val_pred
```

50/50 [=====] - 1s 13ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [False],
               [False]])
```

```
In [ ]: model_name = "NN with SGD"
```

```
train_metric_df.loc[model_name] = recall_score(y_train, y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val, y_val_pred)
```

Classification report

```
In [ ]: #lassification report
cr = classification_report(y_train, y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.80	1.00	0.89	5096
1	1.00	0.00	0.00	1304
accuracy			0.80	6400
macro avg	0.90	0.50	0.44	6400
weighted avg	0.84	0.80	0.71	6400

```
In [ ]: #classification report
cr=classification_report(y_val, y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.80	1.00	0.89	1274
1	1.00	0.00	0.01	326
accuracy			0.80	1600
macro avg	0.90	0.50	0.45	1600
weighted avg	0.84	0.80	0.71	1600

Confusion matrix

```
In [ ]: make_confusion_matrix(y_train, y_train_pred)
```

```
In [ ]: make_confusion_matrix(y_val, y_val_pred)
```



Model Performance Improvement

Neural Network with Adam Optimizer

```
In [ ]: backend.clear_session()
#Fixing the seed for random number generators so that we can ensure we receive the same results
np.random.seed(2)
```

```
random.seed(2)
tf.random.set_seed(2)
```

```
In [ ]: #Initializing the neural network
model_1 = Sequential()
model_1.add(Dense(64,activation='relu',input_dim = X_train.shape[1]))
model_1.add(Dense(32,activation='relu'))
model_1.add(Dense(1, activation = 'sigmoid'))
```

```
In [ ]: # Adam Optimizer
optimizer = tf.keras.optimizers.Adam()
metric = keras.metrics.Recall()
```

```
In [ ]: model_1.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=[metric])
```

```
In [ ]: model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 64)	768
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
=====		
Total params: 2881 (11.25 KB)		
Trainable params: 2881 (11.25 KB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

```
In [ ]: #Fitting the ANN batch_size=32 and epochs=20
history_1 = model_1.fit(
    X_train,y_train,
    batch_size=32,
    validation_data=(X_val,y_val),
    epochs=20,
    verbose=1
)
```

Epoch 1/20
200/200 [=====] - 6s 5ms/step - loss: 0.4473 - recall: 0.1258 - val_loss: 0.4116 - val_recall: 0.2423

Epoch 2/20
200/200 [=====] - 1s 5ms/step - loss: 0.3783 - recall: 0.3658 - val_loss: 0.3767 - val_recall: 0.3957

Epoch 3/20
200/200 [=====] - 2s 12ms/step - loss: 0.3576 - recall: 0.4241 - val_loss: 0.3671 - val_recall: 0.4479

Epoch 4/20
200/200 [=====] - 4s 20ms/step - loss: 0.3491 - recall: 0.4387 - val_loss: 0.3623 - val_recall: 0.4816

Epoch 5/20
200/200 [=====] - 4s 21ms/step - loss: 0.3425 - recall: 0.4578 - val_loss: 0.3604 - val_recall: 0.3620

Epoch 6/20
200/200 [=====] - 4s 20ms/step - loss: 0.3381 - recall: 0.4548 - val_loss: 0.3580 - val_recall: 0.4479

Epoch 7/20
200/200 [=====] - 4s 19ms/step - loss: 0.3358 - recall: 0.4724 - val_loss: 0.3553 - val_recall: 0.4049

Epoch 8/20
200/200 [=====] - 5s 25ms/step - loss: 0.3326 - recall: 0.4594 - val_loss: 0.3560 - val_recall: 0.3988

Epoch 9/20
200/200 [=====] - 5s 23ms/step - loss: 0.3305 - recall: 0.4739 - val_loss: 0.3555 - val_recall: 0.3865

Epoch 10/20
200/200 [=====] - 5s 25ms/step - loss: 0.3280 - recall: 0.4670 - val_loss: 0.3507 - val_recall: 0.5092

Epoch 11/20
200/200 [=====] - 3s 16ms/step - loss: 0.3252 - recall: 0.4824 - val_loss: 0.3567 - val_recall: 0.5429

Epoch 12/20
200/200 [=====] - 4s 19ms/step - loss: 0.3251 - recall: 0.4808 - val_loss: 0.3535 - val_recall: 0.4816

Epoch 13/20
200/200 [=====] - 5s 23ms/step - loss: 0.3229 - recall: 0.4939 - val_loss: 0.3509 - val_recall: 0.4294

Epoch 14/20
200/200 [=====] - 4s 22ms/step - loss: 0.3212 - recall: 0.4847 - val_loss: 0.3539 - val_recall: 0.3896

Epoch 15/20
200/200 [=====] - 4s 19ms/step - loss: 0.3192 - recall: 0.4870 - val_loss: 0.3521 - val_recall: 0.3804

Epoch 16/20
200/200 [=====] - 3s 16ms/step - loss: 0.3188 - recall: 0.4992 - val_loss: 0.3522 - val_recall: 0.4387

Epoch 17/20
200/200 [=====] - 4s 18ms/step - loss: 0.3145 - recall: 0.5038 - val_loss: 0.3486 - val_recall: 0.4540


Epoch 18/20
200/200 [=====] - 4s 21ms/step - loss: 0.3132 - recall: 0.5000 - val_loss: 0.3524 - val_recall: 0.3988

Epoch 19/20
200/200 [=====] - 5s 23ms/step - loss: 0.3124 - recall: 0.5061 - val_loss: 0.3617 - val_recall: 0.3865

Epoch 20/20
200/200 [=====] - 5s 23ms/step - loss: 0.3103 - recall: 0.5015 - val_loss: 0.3491 - val_recall: 0.4693


Loss function

```
In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_1.history['loss'])
plt.plot(history_1.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



Recall

```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_1.history['recall'])
plt.plot(history_1.history['val_recall'])
plt.title('model recall')
plt.ylabel('Recall')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
In [ ]: #Predicting the results using 0.5 as the threshold
y_train_pred = model_1.predict(X_train)
y_train_pred = (y_train_pred > 0.5)
y_train_pred
```

200/200 [=====] - 1s 6ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [False],
               [False]])
```

```
In [ ]: #Predicting the results using 0.5 as the threshold
y_val_pred = model_1.predict(X_val)
y_val_pred = (y_val_pred > 0.5)
y_val_pred
```

50/50 [=====] - 0s 3ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [False],
               [False]])
```

```
In [ ]: model_name = "NN with Adam"
```

```
train_metric_df.loc[model_name] = recall_score(y_train,y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val,y_val_pred)
```

Classification report

```
In [ ]: #lassification report
cr=classification_report(y_train,y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.89	0.96	0.92	5096
1	0.77	0.52	0.62	1304
accuracy			0.87	6400
macro avg	0.83	0.74	0.77	6400
weighted avg	0.86	0.87	0.86	6400

```
In [ ]: #classification report
cr=classification_report(y_val,y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.88	0.96	0.92	1274
1	0.74	0.47	0.58	326
accuracy			0.86	1600
macro avg	0.81	0.71	0.75	1600
weighted avg	0.85	0.86	0.85	1600

Confusion matrix

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_train, y_train_pred)
```



```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_val,y_val_pred)
```



Neural Network with Adam Optimizer and Dropout

```
In [ ]: backend.clear_session()
#Fixing the seed for random number generators so that we can ensure we receive the same results
np.random.seed(2)
random.seed(2)
tf.random.set_seed(2)
```

```
In [ ]: #Initializing the neural network
model_2 = Sequential()
model_2.add(Dense(32,activation='relu',input_dim = X_train.shape[1]))
model_2.add(Dropout(0.2))
model_2.add(Dense(64,activation='relu'))
model_2.add(Dense(32,activation='relu'))
model_2.add(Dropout(0.2))
model_2.add(Dense(16,activation='relu'))
model_2.add(Dense(1, activation = 'sigmoid'))
```

```
In [ ]: # Adam
optimizer = tf.keras.optimizers.Adam()
metric = keras.metrics.Recall()
```

```
In [ ]: model_2.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=[metric])
```

```
In [ ]: # Summary of the model
model_2.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 32)	384

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 32)	384
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 64)	2112
dense_2 (Dense)	(None, 32)	2080
dropout_1 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 16)	528
dense_4 (Dense)	(None, 1)	17

```
=====
Total params: 5121 (20.00 KB)
Trainable params: 5121 (20.00 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
In [ ]: #Fitting the ANN with batch_size = 32 and 100 epochs
history_2 = model_2.fit(
    X_train,y_train,
    batch_size=32,
    epochs=100,
    verbose=1,
    validation_data=(X_val,y_val)
)
```


Epoch 1/100
200/200 [=====] - 6s 17ms/step - loss: 0.4771 - recall: 0.0383 - val_loss: 0.4321 - val_recall: 0.0491

Epoch 2/100
200/200 [=====] - 2s 11ms/step - loss: 0.4246 - recall: 0.1848 - val_loss: 0.4006 - val_recall: 0.1503

Epoch 3/100
200/200 [=====] - 3s 13ms/step - loss: 0.3983 - recall: 0.2845 - val_loss: 0.3756 - val_recall: 0.2577

Epoch 4/100
200/200 [=====] - 2s 8ms/step - loss: 0.3826 - recall: 0.3512 - val_loss: 0.3665 - val_recall: 0.5153

Epoch 5/100
200/200 [=====] - 3s 15ms/step - loss: 0.3766 - recall: 0.3811 - val_loss: 0.3623 - val_recall: 0.3129

Epoch 6/100
200/200 [=====] - 5s 23ms/step - loss: 0.3677 - recall: 0.3773 - val_loss: 0.3603 - val_recall: 0.4233

Epoch 7/100
200/200 [=====] - 4s 19ms/step - loss: 0.3636 - recall: 0.4011 - val_loss: 0.3504 - val_recall: 0.3926

Epoch 8/100
200/200 [=====] - 4s 19ms/step - loss: 0.3587 - recall: 0.4202 - val_loss: 0.3538 - val_recall: 0.3558

Epoch 9/100
200/200 [=====] - 5s 25ms/step - loss: 0.3580 - recall: 0.3965 - val_loss: 0.3494 - val_recall: 0.3681

Epoch 10/100
200/200 [=====] - 5s 24ms/step - loss: 0.3549 - recall: 0.4110 - val_loss: 0.3498 - val_recall: 0.4632

Epoch 11/100
200/200 [=====] - 5s 24ms/step - loss: 0.3486 - recall: 0.4241 - val_loss: 0.3459 - val_recall: 0.4448

Epoch 12/100
200/200 [=====] - 4s 19ms/step - loss: 0.3495 - recall: 0.4248 - val_loss: 0.3558 - val_recall: 0.3834

Epoch 13/100
200/200 [=====] - 5s 27ms/step - loss: 0.3485 - recall: 0.4233 - val_loss: 0.3535 - val_recall: 0.3988

Epoch 14/100
200/200 [=====] - 5s 26ms/step - loss: 0.3495 - recall: 0.4294 - val_loss: 0.3477 - val_recall: 0.3804

Epoch 15/100
200/200 [=====] - 5s 25ms/step - loss: 0.3447 - recall: 0.4348 - val_loss: 0.3446 - val_recall: 0.4233

Epoch 16/100
200/200 [=====] - 4s 21ms/step - loss: 0.3449 - recall: 0.4402 - val_loss: 0.3466 - val_recall: 0.4325

Epoch 17/100
200/200 [=====] - 6s 31ms/step - loss: 0.3492 - recall: 0.4410 - val_loss: 0.3455 - val_recall: 0.3988

Epoch 18/100
200/200 [=====] - 5s 25ms/step - loss: 0.3439 - recall: 0.4586 - val_loss: 0.3435 - val_recall: 0.4264

Epoch 19/100
200/200 [=====] - 6s 28ms/step - loss: 0.3422 - recall: 0.4433 - val_loss: 0.3522 - val_recall: 0.4018

Epoch 20/100
200/200 [=====] - 4s 21ms/step - loss: 0.3419 - recall: 0.4440 - val_loss: 0.3428 - val_recall: 0.4601

Epoch 21/100
200/200 [=====] - 6s 28ms/step - loss: 0.3416 - recall: 0.4509 - val_loss: 0.3471 - val_recall: 0.4325

Epoch 22/100
200/200 [=====] - 5s 26ms/step - loss: 0.3400 - recall: 0.4586 - val_loss: 0.3417 - val_recall: 0.4632

Epoch 23/100
200/200 [=====] - 5s 26ms/step - loss: 0.3384 - recall: 0.4463 - val_loss: 0.3460 - val_recall: 0.4387

Epoch 24/100
200/200 [=====] - 5s 26ms/step - loss: 0.3364 - recall: 0.4624 - val_loss: 0.3452 - val_recall: 0.4601

Epoch 25/100
200/200 [=====] - 6s 31ms/step - loss: 0.3365 - recall: 0.4609 - val_loss: 0.3466 - val_recall: 0.4264

Epoch 26/100
200/200 [=====] - 7s 34ms/step - loss: 0.3367 - recall: 0.4617 - val_loss: 0.3445 - val_recall: 0.4172

Epoch 27/100
200/200 [=====] - 3s 15ms/step - loss: 0.3324 - recall: 0.4732 - val_loss: 0.3479 - val_recall: 0.4325

Epoch 28/100
200/200 [=====] - 2s 12ms/step - loss: 0.3304 - recall: 0.4755 - val_loss: 0.3458 - val_recall: 0.4110

Epoch 29/100
200/200 [=====] - 1s 7ms/step - loss: 0.3364 - recall: 0.4448 - val_loss: 0.3450 - val_recall: 0.4663

Epoch 30/100
200/200 [=====] - 1s 7ms/step - loss: 0.3353 - recall: 0.4655 - val_loss: 0.3469 - val_recall: 0.4202

Epoch 31/100
200/200 [=====] - 2s 9ms/step - loss: 0.3286 - recall: 0.4716 - val_loss: 0.3467 - val_recall: 0.4141

Epoch 32/100
200/200 [=====] - 2s 8ms/step - loss: 0.3342 - recall: 0.4624 - val_loss: 0.3437 - val_recall: 0.4049

Epoch 33/100
200/200 [=====] - 2s 8ms/step - loss: 0.3306 - recall: 0.4632 - val_loss: 0.3428 - val_recall: 0.4601

Epoch 34/100
200/200 [=====] - 2s 8ms/step - loss: 0.3296 - recall: 0.4663 - val_loss: 0.3455 - val_recall: 0.4479

Epoch 35/100
200/200 [=====] - 1s 7ms/step - loss: 0.3307 - recall: 0.4755 - val_loss: 0.3486 - val_recall: 0.4908

Epoch 36/100
200/200 [=====] - 1s 4ms/step - loss: 0.3269 - recall: 0.4755 - val_loss: 0.3432 - val_recall: 0.4571

Epoch 37/100
200/200 [=====] - 2s 10ms/step - loss: 0.3250 - recall: 0.4977 - val_loss: 0.3474 - val_recall: 0.4018

Epoch 38/100
200/200 [=====] - 2s 9ms/step - loss: 0.3244 - recall: 0.4747 - val_loss: 0.3458 - val_recall: 0.4601

Epoch 39/100
200/200 [=====] - 2s 11ms/step - loss: 0.3266 - recall: 0.4831 - val_loss: 0.3484 - val_recall: 0.4663

Epoch 40/100
200/200 [=====] - 4s 21ms/step - loss: 0.3238 - recall: 0.4816 - val_loss: 0.3524 - val_recall: 0.4816

Epoch 41/100
200/200 [=====] - 3s 16ms/step - loss: 0.3228 - recall: 0.4816 - val_loss: 0.3488 - val_recall: 0.4325

Epoch 42/100
200/200 [=====] - 1s 7ms/step - loss: 0.3266 - recall: 0.4755 - val_loss: 0.3456 - val_recall: 0.4202

Epoch 43/100
200/200 [=====] - 1s 4ms/step - loss: 0.3235 - recall: 0.4824 - val_loss: 0.3485 - val_recall: 0.4202

Epoch 44/100
200/200 [=====] - 1s 5ms/step - loss: 0.3217 - recall: 0.4847 - val_loss: 0.3468 - val_recall: 0.4264

Epoch 45/100
200/200 [=====] - 2s 9ms/step - loss: 0.3204 - recall: 0.4908 - val_loss: 0.3470 - val_recall: 0.4417

Epoch 46/100
200/200 [=====] - 1s 5ms/step - loss: 0.3220 - recall: 0.4831 - val_loss: 0.3497 - val_recall: 0.4294

Epoch 47/100
200/200 [=====] - 2s 8ms/step - loss: 0.3133 - recall: 0.4992 - val_loss: 0.3523 - val_recall: 0.4479

Epoch 48/100
200/200 [=====] - 1s 4ms/step - loss: 0.3221 - recall: 0.4793 - val_loss: 0.3506 - val_recall: 0.4110

Epoch 49/100
200/200 [=====] - 1s 6ms/step - loss: 0.3187 - recall: 0.4946 - val_loss: 0.3506 - val_recall: 0.4294

Epoch 50/100
200/200 [=====] - 1s 6ms/step - loss: 0.3206 - recall: 0.4801 - val_loss: 0.3447 - val_recall: 0.4387

Epoch 51/100
200/200 [=====] - 1s 4ms/step - loss: 0.3216 - recall: 0.4985 - val_loss: 0.3461 - val_recall: 0.4110

Epoch 52/100
200/200 [=====] - 1s 4ms/step - loss: 0.3165 - recall: 0.4931 - val_loss: 0.3478 - val_recall: 0.4110

Epoch 53/100
200/200 [=====] - 1s 7ms/step - loss: 0.3153 - recall: 0.5061 - val_loss: 0.3476 - val_recall: 0.4448

Epoch 54/100
200/200 [=====] - 3s 15ms/step - loss: 0.3172 - recall: 0.5077 - val_loss: 0.3508 - val_recall: 0.4049

Epoch 55/100
200/200 [=====] - 4s 22ms/step - loss: 0.3227 - recall: 0.4969 - val_loss: 0.3471 - val_recall: 0.4018

Epoch 56/100
200/200 [=====] - 4s 22ms/step - loss: 0.3184 - recall: 0.5023 - val_loss: 0.3480 - val_recall: 0.4264

Epoch 57/100
200/200 [=====] - 4s 21ms/step - loss: 0.3167 - recall: 0.4939 - val_loss: 0.3462 - val_recall: 0.4325

Epoch 58/100
200/200 [=====] - 4s 21ms/step - loss: 0.3150 - recall: 0.5153 - val_loss: 0.3492 - val_recall: 0.4356

Epoch 59/100
200/200 [=====] - 4s 22ms/step - loss: 0.3174 - recall: 0.4939 - val_loss: 0.3473 - val_recall: 0.4540

Epoch 60/100
200/200 [=====] - 4s 21ms/step - loss: 0.3177 - recall: 0.5123 - val_loss: 0.3502 - val_recall: 0.3804

Epoch 61/100
200/200 [=====] - 5s 25ms/step - loss: 0.3131 - recall: 0.5092 - val_loss: 0.3523 - val_recall: 0.4294

Epoch 62/100
200/200 [=====] - 4s 22ms/step - loss: 0.3141 - recall: 0.5046 - val_loss: 0.3479 - val_recall: 0.4571

Epoch 63/100
200/200 [=====] - 5s 26ms/step - loss: 0.3130 - recall: 0.5184 - val_loss: 0.3481 - val_recall: 0.4693

Epoch 64/100
200/200 [=====] - 5s 23ms/step - loss: 0.3154 - recall: 0.5169 - val_loss: 0.3462 - val_recall: 0.4387

Epoch 65/100
200/200 [=====] - 5s 26ms/step - loss: 0.3155 - recall: 0.5046 - val_loss: 0.3508 - val_recall: 0.3834

Epoch 66/100
200/200 [=====] - 5s 23ms/step - loss: 0.3124 - recall: 0.5123 - val_loss: 0.3464 - val_recall: 0.4448

Epoch 67/100
200/200 [=====] - 5s 24ms/step - loss: 0.3085 - recall: 0.5222 - val_loss: 0.3502 - val_recall: 0.5092

Epoch 68/100
200/200 [=====] - 5s 25ms/step - loss: 0.3138 - recall: 0.4962 - val_loss: 0.3518 - val_recall: 0.4356

Epoch 69/100
200/200 [=====] - 4s 21ms/step - loss: 0.3111 - recall: 0.5100 - val_loss: 0.3540 - val_recall: 0.4387

Epoch 70/100
200/200 [=====] - 5s 24ms/step - loss: 0.3110 - recall: 0.5184 - val_loss: 0.3513 - val_recall: 0.4571

Epoch 71/100
200/200 [=====] - 5s 24ms/step - loss: 0.3157 - recall: 0.5000 - val_loss: 0.3492 - val_recall: 0.4233

Epoch 72/100
200/200 [=====] - 5s 25ms/step - loss: 0.3104 - recall: 0.5176 - val_loss: 0.3506 - val_recall: 0.4479

Epoch 73/100
200/200 [=====] - 6s 29ms/step - loss: 0.3157 - recall: 0.5069 - val_loss: 0.3459 - val_recall: 0.4540

Epoch 74/100
200/200 [=====] - 4s 20ms/step - loss: 0.3091 - recall: 0.5176 - val_loss: 0.3514 - val_recall: 0.4571

Epoch 75/100
200/200 [=====] - 2s 10ms/step - loss: 0.3107 - recall: 0.5184 - val_loss: 0.3520 - val_recall: 0.3926

Epoch 76/100
200/200 [=====] - 3s 15ms/step - loss: 0.3080 - recall: 0.4977 - val_loss: 0.3474 - val_recall: 0.4325

Epoch 77/100
200/200 [=====] - 5s 25ms/step - loss: 0.3028 - recall: 0.5215 - val_loss: 0.3525 - val_recall: 0.4264

Epoch 78/100
200/200 [=====] - 4s 20ms/step - loss: 0.3058 - recall: 0.5123 - val_loss: 0.3480 - val_recall: 0.4049

Epoch 79/100
200/200 [=====] - 5s 25ms/step - loss: 0.3087 - recall: 0.5061 - val_loss: 0.3500 - val_recall: 0.4294

Epoch 80/100
200/200 [=====] - 4s 20ms/step - loss: 0.3106 - recall: 0.5107 - val_loss: 0.3533 - val_recall: 0.4141

Epoch 81/100
200/200 [=====] - 4s 20ms/step - loss: 0.3099 - recall: 0.5084 - val_loss: 0.3506 - val_recall: 0.4264

Epoch 82/100
200/200 [=====] - 4s 22ms/step - loss: 0.3058 - recall: 0.5230 - val_loss: 0.3508 - val_recall: 0.4632

Epoch 83/100
200/200 [=====] - 5s 24ms/step - loss: 0.3014 - recall: 0.5261 - val_loss: 0.3500 - val_recall: 0.4356

Epoch 84/100
200/200 [=====] - 5s 26ms/step - loss: 0.3042 - recall: 0.5199 - val_loss: 0.3606 - val_recall: 0.4571

Epoch 85/100
200/200 [=====] - 5s 23ms/step - loss: 0.3064 - recall: 0.5169 - val_loss: 0.3484 - val_recall: 0.4110

Epoch 86/100
200/200 [=====] - 6s 31ms/step - loss: 0.3037 - recall: 0.5199 - val_loss: 0.3559 - val_recall: 0.4693

Epoch 87/100
200/200 [=====] - 3s 15ms/step - loss: 0.3042 - recall: 0.5276 - val_loss: 0.3571 - val_recall: 0.4571

Epoch 88/100
200/200 [=====] - 4s 20ms/step - loss: 0.3082 - recall: 0.5207 - val_loss: 0.3503 - val_recall: 0.4755

Epoch 89/100
200/200 [=====] - 3s 17ms/step - loss: 0.3091 - recall: 0.5238 - val_loss: 0.3503 - val_recall: 0.4202

Epoch 90/100
200/200 [=====] - 4s 19ms/step - loss: 0.3051 - recall: 0.5192 - val_loss: 0.3510 - val_recall: 0.4540

Epoch 91/100
200/200 [=====] - 5s 23ms/step - loss: 0.3009 - recall: 0.5445 - val_loss: 0.3577 - val_recall: 0.4018

Epoch 92/100
200/200 [=====] - 5s 23ms/step - loss: 0.3048 - recall: 0.5337 - val_loss: 0.3493 - val_recall: 0.4356

Epoch 93/100
200/200 [=====] - 4s 22ms/step - loss: 0.3062 - recall: 0.5084 - val_loss: 0.3518 - val_recall: 0.4264

Epoch 94/100
200/200 [=====] - 3s 17ms/step - loss: 0.3051 - recall: 0.5291 - val_loss: 0.3504 - val_recall: 0.4969

Epoch 95/100
200/200 [=====] - 4s 22ms/step - loss: 0.3055 - recall: 0.5337 - val_loss: 0.3534 - val_recall: 0.4417

Epoch 96/100
200/200 [=====] - 6s 29ms/step - loss: 0.3029 - recall: 0.5399 - val_loss: 0.3564 - val_recall: 0.4264

Epoch 97/100
200/200 [=====] - 5s 22ms/step - loss: 0.3042 - recall: 0.5261 - val_loss: 0.3553 - val_recall: 0.4387


Epoch 98/100
200/200 [=====] - 5s 24ms/step - loss: 0.3019 - recall: 0.5261 - val_loss: 0.3575 - val_recall: 0.4571

Epoch 99/100
200/200 [=====] - 5s 24ms/step - loss: 0.3061 - recall: 0.5192 - val_loss: 0.3548 - val_recall: 0.4540


Epoch 100/100
200/200 [=====] - 5s 27ms/step - loss: 0.2994 - recall: 0.5391 - val_loss: 0.3586 - val_recall: 0.4908

Loss function

```
In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_2.history['loss'])
plt.plot(history_2.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_2.history['recall'])
plt.plot(history_2.history['val_recall'])
plt.title('model recall')
plt.ylabel('recall')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
In [ ]: #Predicting the results using best as a threshold
y_train_pred = model_2.predict(X_train)
y_train_pred = (y_train_pred > 0.5)
y_train_pred
```

200/200 [=====] - 4s 10ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [ True],
               [False]])
```

```
In [ ]: #Predicting the results using 0.5 as the threshold.
y_val_pred = model_2.predict(X_val)
y_val_pred = (y_val_pred > 0.5)
y_val_pred
```

50/50 [=====] - 1s 7ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [ True],
               [ True]])
```

```
In [ ]: model_name = "NN with Adam & Dropout"
```

```
train_metric_df.loc[model_name] = recall_score(y_train,y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val,y_val_pred)
```

Classification report

```
In [ ]: #classification report
cr=classification_report(y_train,y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.90	0.97	0.93	5096
1	0.82	0.60	0.69	1304
accuracy			0.89	6400
macro avg	0.86	0.78	0.81	6400
weighted avg	0.89	0.89	0.89	6400

```
In [ ]: #classification report
cr = classification_report(y_val,y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.88	0.94	0.91	1274
1	0.68	0.49	0.57	326
accuracy			0.85	1600
macro avg	0.78	0.72	0.74	1600
weighted avg	0.84	0.85	0.84	1600

Confusion matrix

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_train, y_train_pred)
```

```
In [ ]: #Calculating the confusion matrix  
make_confusion_matrix(y_val,y_val_pred)
```



Neural Network with Balanced Data (by applying SMOTE) and SGD Optimizer

Let's try to apply SMOTE to balance this dataset and then again apply hyperparameter tuning accordingly.

```
In [ ]: sm = SMOTE(random_state=42)
#Complete the code to fit SMOTE on the training data.
X_train_smote, y_train_smote= sm.fit_resample(X_train,y_train)
print('After UpSampling, the shape of train_X: {}'.format(X_train_smote.shape))
print('After UpSampling, the shape of train_y: {} \n'.format(y_train_smote.shape))
```

After UpSampling, the shape of train_X: (10192, 11)
 After UpSampling, the shape of train_y: (10192,)

Let's build a model with the balanced dataset

```
In [ ]: backend.clear_session()
#Fixing the seed for random number generators so that we can ensure we receive the same
np.random.seed(2)
random.seed(2)
tf.random.set_seed(2)
```

```
In [ ]: #Initializing the model
model_3 = Sequential()
model_3.add(Dense(32,activation='relu',input_dim = X_train_smote.shape[1]))
model_3.add(Dense(16,activation='relu'))
model_3.add(Dense(8,activation='relu'))
model_3.add(Dense(1, activation = 'sigmoid'))
```

```
In [ ]: # SGD
optimizer = tf.keras.optimizers.SGD(0.001)
metric = keras.metrics.Recall()
```

```
In [ ]: model_3.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=[metric])
```

```
In [ ]: model_3.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 32)	384
=====		
Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 32)	384
dense_1 (Dense)	(None, 16)	528
dense_2 (Dense)	(None, 8)	136
dense_3 (Dense)	(None, 1)	9

=====

Total params: 1057 (4.13 KB)
 Trainable params: 1057 (4.13 KB)
 Non-trainable params: 0 (0.00 Byte)

=====

```
In [ ]: #Fitting the ANN
history_3 = model_3.fit(
    X_train_smote, y_train_smote,
```

```
batch_size=32,  
epochs=50,  
verbose=1,  
validation_data = (X_val,y_val)  
)
```

Epoch 1/50
319/319 [=====] - 15s 37ms/step - loss: 0.6995 - recall: 0.7983 - val_loss: 0.7144 - val_recall: 0.8067

Epoch 2/50
319/319 [=====] - 5s 16ms/step - loss: 0.6961 - recall: 0.7855 - val_loss: 0.7066 - val_recall: 0.7822

Epoch 3/50
319/319 [=====] - 7s 23ms/step - loss: 0.6927 - recall: 0.7628 - val_loss: 0.6988 - val_recall: 0.7362

Epoch 4/50
319/319 [=====] - 6s 20ms/step - loss: 0.6892 - recall: 0.7223 - val_loss: 0.6909 - val_recall: 0.6902

Epoch 5/50
319/319 [=====] - 5s 17ms/step - loss: 0.6854 - recall: 0.6811 - val_loss: 0.6830 - val_recall: 0.6626

Epoch 6/50
319/319 [=====] - 7s 21ms/step - loss: 0.6816 - recall: 0.6599 - val_loss: 0.6755 - val_recall: 0.6472

Epoch 7/50
319/319 [=====] - 3s 10ms/step - loss: 0.6779 - recall: 0.6348 - val_loss: 0.6688 - val_recall: 0.6258

Epoch 8/50
319/319 [=====] - 4s 11ms/step - loss: 0.6744 - recall: 0.6307 - val_loss: 0.6627 - val_recall: 0.6227

Epoch 9/50
319/319 [=====] - 3s 9ms/step - loss: 0.6710 - recall: 0.6279 - val_loss: 0.6570 - val_recall: 0.6135

Epoch 10/50
319/319 [=====] - 3s 8ms/step - loss: 0.6677 - recall: 0.6305 - val_loss: 0.6518 - val_recall: 0.6104

Epoch 11/50
319/319 [=====] - 4s 11ms/step - loss: 0.6644 - recall: 0.6315 - val_loss: 0.6469 - val_recall: 0.6043

Epoch 12/50
319/319 [=====] - 4s 12ms/step - loss: 0.6611 - recall: 0.6391 - val_loss: 0.6419 - val_recall: 0.5920

Epoch 13/50
319/319 [=====] - 3s 10ms/step - loss: 0.6578 - recall: 0.6336 - val_loss: 0.6373 - val_recall: 0.5890

Epoch 14/50
319/319 [=====] - 4s 14ms/step - loss: 0.6544 - recall: 0.6415 - val_loss: 0.6327 - val_recall: 0.5951

Epoch 15/50
319/319 [=====] - 2s 7ms/step - loss: 0.6509 - recall: 0.6419 - val_loss: 0.6283 - val_recall: 0.5982

Epoch 16/50
319/319 [=====] - 5s 14ms/step - loss: 0.6474 - recall: 0.6468 - val_loss: 0.6241 - val_recall: 0.6074

Epoch 17/50
319/319 [=====] - 5s 14ms/step - loss: 0.6439 - recall: 0.6531 - val_loss: 0.6200 - val_recall: 0.6074

Epoch 18/50
319/319 [=====] - 3s 10ms/step - loss: 0.6402 - recall: 0.6599 - val_loss: 0.6159 - val_recall: 0.6104

Epoch 19/50
319/319 [=====] - 4s 13ms/step - loss: 0.6365 - recall: 0.6635 - val_loss: 0.6119 - val_recall: 0.6104

Epoch 20/50
319/319 [=====] - 4s 13ms/step - loss: 0.6328 - recall: 0.6644 - val_loss: 0.6083 - val_recall: 0.6104

Epoch 21/50
319/319 [=====] - 5s 17ms/step - loss: 0.6290 - recall: 0.6745 - val_loss: 0.6048 - val_recall: 0.6074

Epoch 22/50
319/319 [=====] - 3s 9ms/step - loss: 0.6251 - recall: 0.6829 - val_loss: 0.6011 - val_recall: 0.6074

Epoch 23/50
319/319 [=====] - 3s 8ms/step - loss: 0.6211 - recall: 0.6849 - val_loss: 0.5979 - val_recall: 0.6074

Epoch 24/50
319/319 [=====] - 3s 8ms/step - loss: 0.6171 - recall: 0.6913 - val_loss: 0.5947 - val_recall: 0.6104

Epoch 25/50
319/319 [=====] - 3s 11ms/step - loss: 0.6131 - recall: 0.6931 - val_loss: 0.5920 - val_recall: 0.6135

Epoch 26/50
319/319 [=====] - 3s 9ms/step - loss: 0.6090 - recall: 0.6980 - val_loss: 0.5891 - val_recall: 0.6227

Epoch 27/50
319/319 [=====] - 3s 9ms/step - loss: 0.6049 - recall: 0.7019 - val_loss: 0.5863 - val_recall: 0.6319

Epoch 28/50
319/319 [=====] - 2s 6ms/step - loss: 0.6007 - recall: 0.7082 - val_loss: 0.5835 - val_recall: 0.6350

Epoch 29/50
319/319 [=====] - 2s 7ms/step - loss: 0.5966 - recall: 0.7113 - val_loss: 0.5812 - val_recall: 0.6472

Epoch 30/50
319/319 [=====] - 3s 11ms/step - loss: 0.5925 - recall: 0.7217 - val_loss: 0.5781 - val_recall: 0.6472

Epoch 31/50
319/319 [=====] - 4s 12ms/step - loss: 0.5884 - recall: 0.7237 - val_loss: 0.5757 - val_recall: 0.6595

Epoch 32/50
319/319 [=====] - 3s 10ms/step - loss: 0.5843 - recall: 0.7280 - val_loss: 0.5730 - val_recall: 0.6626

Epoch 33/50
319/319 [=====] - 4s 11ms/step - loss: 0.5803 - recall: 0.7310 - val_loss: 0.5705 - val_recall: 0.6748

Epoch 34/50
319/319 [=====] - 4s 12ms/step - loss: 0.5763 - recall: 0.7351 - val_loss: 0.5678 - val_recall: 0.6748

Epoch 35/50
319/319 [=====] - 4s 11ms/step - loss: 0.5724 - recall: 0.7384 - val_loss: 0.5648 - val_recall: 0.6779

Epoch 36/50
319/319 [=====] - 4s 12ms/step - loss: 0.5686 - recall: 0.7370 - val_loss: 0.5626 - val_recall: 0.6810

Epoch 37/50
319/319 [=====] - 3s 10ms/step - loss: 0.5649 - recall: 0.7398 - val_loss: 0.5613 - val_recall: 0.6840

Epoch 38/50
319/319 [=====] - 4s 12ms/step - loss: 0.5614 - recall: 0.7457 - val_loss: 0.5588 - val_recall: 0.6871

Epoch 39/50
319/319 [=====] - 3s 10ms/step - loss: 0.5580 - recall: 0.7443 - val_loss: 0.5576 - val_recall: 0.6902

Epoch 40/50
319/319 [=====] - 3s 10ms/step - loss: 0.5547 - recall: 0.7490 - val_loss: 0.5553 - val_recall: 0.6902

```

Epoch 41/50
319/319 [=====] - 3s 10ms/step - loss: 0.5516 - recall: 0.75
24 - val_loss: 0.5531 - val_recall: 0.6871
Epoch 42/50
319/319 [=====] - 3s 10ms/step - loss: 0.5486 - recall: 0.75
29 - val_loss: 0.5509 - val_recall: 0.6902
Epoch 43/50
319/319 [=====] - 3s 10ms/step - loss: 0.5458 - recall: 0.75
27 - val_loss: 0.5499 - val_recall: 0.6994
Epoch 44/50
319/319 [=====] - 3s 9ms/step - loss: 0.5430 - recall: 0.754
9 - val_loss: 0.5486 - val_recall: 0.7025
Epoch 45/50
319/319 [=====] - 2s 7ms/step - loss: 0.5404 - recall: 0.755
7 - val_loss: 0.5473 - val_recall: 0.7086
Epoch 46/50
319/319 [=====] - 4s 12ms/step - loss: 0.5379 - recall: 0.76
08 - val_loss: 0.5461 - val_recall: 0.7086
Epoch 47/50
319/319 [=====] - 4s 13ms/step - loss: 0.5354 - recall: 0.75
84 - val_loss: 0.5463 - val_recall: 0.7147
Epoch 48/50
319/319 [=====] - 3s 9ms/step - loss: 0.5331 - recall: 0.763
1 - val_loss: 0.5443 - val_recall: 0.7147
Epoch 49/50
319/319 [=====] - 4s 13ms/step - loss: 0.5308 - recall: 0.76
28 - val_loss: 0.5432 - val_recall: 0.7147
Epoch 50/50
319/319 [=====] - 4s 12ms/step - loss: 0.5286 - recall: 0.76
45 - val_loss: 0.5425 - val_recall: 0.7147


```

Loss function


```

In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_3.history['loss'])
plt.plot(history_3.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

```



```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_3.history['recall'])
plt.plot(history_3.history['val_recall'])
plt.title('model recall')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
In [ ]: y_train_pred = model_3.predict(X_train_smote)
        #Predicting the results using 0.5 as the threshold
        y_train_pred = (y_train_pred > 0.5)
        y_train_pred
```

319/319 [=====] - 4s 11ms/step

```
Out[ ]: array([[ True],
               [False],
               [False],
               ...,
               [ True],
               [ True],
               [ True]])
```

```
In [ ]: y_val_pred = model_3.predict(X_val)
        #Predicting the results using 0.5 as the threshold
        y_val_pred = (y_val_pred > 0.5)
        y_val_pred
```

50/50 [=====] - 1s 14ms/step

```
Out[ ]: array([[ True],
               [False],
               [False],
               ...,
               [False],
               [ True],
               [ True]])
```

```
In [ ]: model_name = "NN with SMOTE & SGD"
```



```
train_metric_df.loc[model_name] = recall_score(y_train_smote,y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val,y_val_pred)
```

Classification report

```
In [ ]: cr=classification_report(y_train_smote,y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.76	0.73	0.74	5096
1	0.74	0.77	0.75	5096
accuracy			0.75	10192
macro avg	0.75	0.75	0.75	10192
weighted avg	0.75	0.75	0.75	10192

```
In [ ]: cr=classification_report(y_val,y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.91	0.72	0.81	1274
1	0.40	0.71	0.51	326
accuracy			0.72	1600
macro avg	0.65	0.72	0.66	1600
weighted avg	0.80	0.72	0.75	1600

Confusion matrix

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_train_smote, y_train_pred)
```



```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_val,y_val_pred)
```



Neural Network with Balanced Data (by applying SMOTE) and Adam Optimizer

Let's build a model with the balanced dataset

```
In [ ]: backend.clear_session()
#Fixing the seed for random number generators so that we can ensure we receive the same results
np.random.seed(2)
random.seed(2)
tf.random.set_seed(2)
```

```
In [ ]: #Initializing the model
model_4 = Sequential()
model_4.add(Dense(64,activation='relu',input_dim = X_train_smote.shape[1]))
model_4.add(Dense(32,activation='relu'))
model_4.add(Dense(16,activation='relu'))
model_4.add(Dense(1, activation = 'sigmoid'))
```

```
In [ ]: model_4.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	768
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 16)	528
dense (Dense)	(None, 64)	768
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 16)	528
dense_3 (Dense)	(None, 1)	17

=====
 Total params: 3393 (13.25 KB)
 Trainable params: 3393 (13.25 KB)
 Non-trainable params: 0 (0.00 Byte)

```
In [ ]: # Adam
optimizer = tf.keras.optimizers.Adam()
metric = keras.metrics.Recall()
```

```
In [ ]: model_4.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=[metric])
```

```
In [ ]: model_4.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	768
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 16)	528
dense_3 (Dense)	(None, 1)	17

=====
 Total params: 3393 (13.25 KB)
 Trainable params: 3393 (13.25 KB)
 Non-trainable params: 0 (0.00 Byte)

```
In [ ]: #Fitting the ANN
history_4 = model_4.fit(
    X_train_smote,y_train_smote,
    batch_size=32,
```

```
epochs=50,  
verbose=1,  
validation_data = (X_val,y_val)  
)
```

Epoch 1/50
319/319 [=====] - 3s 4ms/step - loss: 0.5139 - recall: 0.7549 - val_loss: 0.5364 - val_recall: 0.8006

Epoch 2/50
319/319 [=====] - 1s 3ms/step - loss: 0.4405 - recall: 0.7942 - val_loss: 0.4504 - val_recall: 0.7117

Epoch 3/50
319/319 [=====] - 1s 3ms/step - loss: 0.4239 - recall: 0.7957 - val_loss: 0.4842 - val_recall: 0.7515

Epoch 4/50
319/319 [=====] - 1s 4ms/step - loss: 0.4131 - recall: 0.7995 - val_loss: 0.4363 - val_recall: 0.6810

Epoch 5/50
319/319 [=====] - 1s 3ms/step - loss: 0.4048 - recall: 0.8020 - val_loss: 0.5027 - val_recall: 0.7638

Epoch 6/50
319/319 [=====] - 1s 3ms/step - loss: 0.3975 - recall: 0.8138 - val_loss: 0.4654 - val_recall: 0.7025

Epoch 7/50
319/319 [=====] - 1s 5ms/step - loss: 0.3911 - recall: 0.8175 - val_loss: 0.4358 - val_recall: 0.6687

Epoch 8/50
319/319 [=====] - 1s 4ms/step - loss: 0.3833 - recall: 0.8187 - val_loss: 0.4289 - val_recall: 0.6718

Epoch 9/50
319/319 [=====] - 1s 4ms/step - loss: 0.3766 - recall: 0.8308 - val_loss: 0.4211 - val_recall: 0.6595

Epoch 10/50
319/319 [=====] - 1s 4ms/step - loss: 0.3736 - recall: 0.8287 - val_loss: 0.4278 - val_recall: 0.6595

Epoch 11/50
319/319 [=====] - 1s 4ms/step - loss: 0.3686 - recall: 0.8279 - val_loss: 0.4478 - val_recall: 0.6963

Epoch 12/50
319/319 [=====] - 1s 4ms/step - loss: 0.3618 - recall: 0.8338 - val_loss: 0.4205 - val_recall: 0.6166

Epoch 13/50
319/319 [=====] - 2s 5ms/step - loss: 0.3554 - recall: 0.8483 - val_loss: 0.4727 - val_recall: 0.7331

Epoch 14/50
319/319 [=====] - 1s 4ms/step - loss: 0.3505 - recall: 0.8491 - val_loss: 0.4727 - val_recall: 0.6933

Epoch 15/50
319/319 [=====] - 1s 3ms/step - loss: 0.3448 - recall: 0.8518 - val_loss: 0.4510 - val_recall: 0.6810

Epoch 16/50
319/319 [=====] - 1s 4ms/step - loss: 0.3381 - recall: 0.8526 - val_loss: 0.4627 - val_recall: 0.7147

Epoch 17/50
319/319 [=====] - 1s 4ms/step - loss: 0.3364 - recall: 0.8564 - val_loss: 0.5060 - val_recall: 0.7362

Epoch 18/50
319/319 [=====] - 1s 3ms/step - loss: 0.3301 - recall: 0.8620 - val_loss: 0.4618 - val_recall: 0.6779

Epoch 19/50
319/319 [=====] - 1s 3ms/step - loss: 0.3258 - recall: 0.8620 - val_loss: 0.4359 - val_recall: 0.6043

Epoch 20/50
319/319 [=====] - 1s 3ms/step - loss: 0.3200 - recall: 0.8683 - val_loss: 0.4780 - val_recall: 0.6994

Epoch 21/50
319/319 [=====] - 1s 3ms/step - loss: 0.3154 - recall: 0.869
9 - val_loss: 0.4477 - val_recall: 0.5798

Epoch 22/50
319/319 [=====] - 1s 3ms/step - loss: 0.3127 - recall: 0.872
1 - val_loss: 0.4771 - val_recall: 0.6902

Epoch 23/50
319/319 [=====] - 1s 3ms/step - loss: 0.3053 - recall: 0.879
7 - val_loss: 0.4518 - val_recall: 0.6196

Epoch 24/50
319/319 [=====] - 1s 3ms/step - loss: 0.3006 - recall: 0.880
1 - val_loss: 0.4950 - val_recall: 0.6718

Epoch 25/50
319/319 [=====] - 1s 4ms/step - loss: 0.2980 - recall: 0.885
4 - val_loss: 0.4914 - val_recall: 0.6442

Epoch 26/50
319/319 [=====] - 1s 4ms/step - loss: 0.2919 - recall: 0.888
1 - val_loss: 0.5009 - val_recall: 0.6595

Epoch 27/50
319/319 [=====] - 1s 3ms/step - loss: 0.2914 - recall: 0.885
0 - val_loss: 0.4877 - val_recall: 0.6718

Epoch 28/50
319/319 [=====] - 1s 3ms/step - loss: 0.2851 - recall: 0.889
3 - val_loss: 0.4641 - val_recall: 0.5767

Epoch 29/50
319/319 [=====] - 1s 3ms/step - loss: 0.2857 - recall: 0.889
7 - val_loss: 0.4864 - val_recall: 0.6380

Epoch 30/50
319/319 [=====] - 1s 3ms/step - loss: 0.2784 - recall: 0.897
4 - val_loss: 0.5500 - val_recall: 0.6933

Epoch 31/50
319/319 [=====] - 1s 3ms/step - loss: 0.2768 - recall: 0.896
0 - val_loss: 0.4990 - val_recall: 0.6442

Epoch 32/50
319/319 [=====] - 1s 4ms/step - loss: 0.2705 - recall: 0.902
1 - val_loss: 0.4933 - val_recall: 0.5644

Epoch 33/50
319/319 [=====] - 1s 3ms/step - loss: 0.2705 - recall: 0.899
5 - val_loss: 0.5369 - val_recall: 0.6380

Epoch 34/50
319/319 [=====] - 1s 3ms/step - loss: 0.2690 - recall: 0.901
7 - val_loss: 0.4861 - val_recall: 0.5828

Epoch 35/50
319/319 [=====] - 1s 3ms/step - loss: 0.2647 - recall: 0.902
7 - val_loss: 0.5126 - val_recall: 0.6196

Epoch 36/50
319/319 [=====] - 1s 3ms/step - loss: 0.2609 - recall: 0.900
5 - val_loss: 0.4999 - val_recall: 0.6166

Epoch 37/50
319/319 [=====] - 2s 5ms/step - loss: 0.2607 - recall: 0.902
5 - val_loss: 0.5507 - val_recall: 0.6871

Epoch 38/50
319/319 [=====] - 1s 4ms/step - loss: 0.2543 - recall: 0.910
3 - val_loss: 0.5308 - val_recall: 0.6196

Epoch 39/50
319/319 [=====] - 1s 3ms/step - loss: 0.2514 - recall: 0.908
4 - val_loss: 0.5532 - val_recall: 0.6656

Epoch 40/50
319/319 [=====] - 1s 3ms/step - loss: 0.2504 - recall: 0.911
7 - val_loss: 0.5710 - val_recall: 0.6595

```

Epoch 41/50
319/319 [=====] - 1s 3ms/step - loss: 0.2458 - recall: 0.912
7 - val_loss: 0.5389 - val_recall: 0.6319
Epoch 42/50
319/319 [=====] - 1s 3ms/step - loss: 0.2454 - recall: 0.907
4 - val_loss: 0.5229 - val_recall: 0.5368
Epoch 43/50
319/319 [=====] - 1s 3ms/step - loss: 0.2470 - recall: 0.912
1 - val_loss: 0.5356 - val_recall: 0.6135
Epoch 44/50
319/319 [=====] - 1s 3ms/step - loss: 0.2398 - recall: 0.916
0 - val_loss: 0.5442 - val_recall: 0.5828
Epoch 45/50
319/319 [=====] - 1s 3ms/step - loss: 0.2371 - recall: 0.916
8 - val_loss: 0.5312 - val_recall: 0.5890
Epoch 46/50
319/319 [=====] - 1s 3ms/step - loss: 0.2361 - recall: 0.916
8 - val_loss: 0.5648 - val_recall: 0.6472
Epoch 47/50
319/319 [=====] - 1s 3ms/step - loss: 0.2338 - recall: 0.919
3 - val_loss: 0.7073 - val_recall: 0.7117
Epoch 48/50
319/319 [=====] - 1s 3ms/step - loss: 0.2341 - recall: 0.917
8 - val_loss: 0.5834 - val_recall: 0.6411
Epoch 49/50
319/319 [=====] - 1s 3ms/step - loss: 0.2278 - recall: 0.921
1 - val_loss: 0.5970 - val_recall: 0.6319
Epoch 50/50
319/319 [=====] - 1s 3ms/step - loss: 0.2265 - recall: 0.922
3 - val_loss: 0.5785 - val_recall: 0.6043


```

Loss function

```

In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_4.history['loss'])
plt.plot(history_4.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

```



```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_4.history['recall'])
plt.plot(history_4.history['val_recall'])
plt.title('model recall')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```




```
In [ ]: y_train_pred = model_4.predict(X_train_smote)
#Predicting the results using 0.5 as the threshold
y_train_pred = (y_train_pred > 0.5)
y_train_pred
```

319/319 [=====] - 1s 2ms/step

```
Out[ ]: array([[ True],
 [False],
 [False],
 ...,
 [ True],
 [ True],
 [ True]])
```

```
In [ ]: y_val_pred = model_4.predict(X_val)
#Predicting the results using 0.5 as the threshold
y_val_pred = (y_val_pred > 0.5)
y_val_pred
```

50/50 [=====] - 0s 2ms/step

```
Out[ ]: array([[ True],
 [False],
 [False],
 ...,
 [False],
 [ True],
 [ True]])
```

```
In [ ]: model_name = "NN with SMOTE & Adam"
```

```
train_metric_df.loc[model_name] = recall_score(y_train_smote,y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val,y_val_pred)
```

Classification report

```
In [ ]: cr=classification_report(y_train_smote,y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.94	0.90	0.92	5096
1	0.90	0.94	0.92	5096
accuracy			0.92	10192
macro avg	0.92	0.92	0.92	10192
weighted avg	0.92	0.92	0.92	10192

```
In [ ]: cr=classification_report(y_val,y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.89	0.84	0.86	1274
1	0.48	0.60	0.54	326
accuracy			0.79	1600
macro avg	0.69	0.72	0.70	1600
weighted avg	0.81	0.79	0.80	1600

Confusion matrix

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_train_smote, y_train_pred)
```



```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_val,y_val_pred)
```



Neural Network with Balanced Data (by applying SMOTE), Adam Optimizer, and Dropout

```
In [ ]: backend.clear_session()
#Fixing the seed for random number generators so that we can ensure we receive the same results
np.random.seed(2)
random.seed(2)
tf.random.set_seed(2)
```

```
In [ ]: # Initializing the model
model_5 = Sequential()
# Adding the input layer
model_5.add(Dense(64, activation='relu', input_dim=X_train_smote.shape[1]))
# Adding dropout regularization
model_5.add(Dropout(0.5)) # Dropout rate of 0.5 (50% dropout)
# Adding a hidden layer
model_5.add(Dense(32, activation='relu'))
# Adding dropout regularization
model_5.add(Dropout(0.3)) # Dropout rate of 0.3 (30% dropout)
# Adding another hidden layer
model_5.add(Dense(8, activation='relu'))
# Adding the output layer
model_5.add(Dense(1, activation='sigmoid'))
```

```
In [ ]: # Adam
optimizer = tf.keras.optimizers.Adam()
metric = keras.metrics.Recall()
```

```
In [ ]: model_5.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=[metric])
```

```
In [ ]: model_5.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 64)	768
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dropout_1 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 8)	264
dense_3 (Dense)	(None, 1)	9
=====		
Total params: 3121 (12.19 KB)		
Trainable params: 3121 (12.19 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [ ]: history_5 = model_5.fit(
        X_train_smote,y_train_smote,
        batch_size=32,
        epochs=50,
        verbose=1,
        validation_data = (X_val,y_val))
```

Epoch 1/50
319/319 [=====] - 7s 10ms/step - loss: 0.6190 - recall: 0.6782 - val_loss: 0.5435 - val_recall: 0.7025

Epoch 2/50
319/319 [=====] - 1s 3ms/step - loss: 0.5411 - recall: 0.7353 - val_loss: 0.4899 - val_recall: 0.7117

Epoch 3/50
319/319 [=====] - 5s 16ms/step - loss: 0.5112 - recall: 0.7575 - val_loss: 0.4611 - val_recall: 0.6933

Epoch 4/50
319/319 [=====] - 3s 10ms/step - loss: 0.4919 - recall: 0.7594 - val_loss: 0.4611 - val_recall: 0.6748

Epoch 5/50
319/319 [=====] - 3s 9ms/step - loss: 0.4787 - recall: 0.7531 - val_loss: 0.4586 - val_recall: 0.6779

Epoch 6/50
319/319 [=====] - 3s 10ms/step - loss: 0.4756 - recall: 0.7563 - val_loss: 0.4474 - val_recall: 0.6963

Epoch 7/50
319/319 [=====] - 3s 11ms/step - loss: 0.4676 - recall: 0.7610 - val_loss: 0.4448 - val_recall: 0.6748

Epoch 8/50
319/319 [=====] - 4s 11ms/step - loss: 0.4661 - recall: 0.7649 - val_loss: 0.4284 - val_recall: 0.6810

Epoch 9/50
319/319 [=====] - 3s 11ms/step - loss: 0.4605 - recall: 0.7730 - val_loss: 0.4577 - val_recall: 0.7147

Epoch 10/50
319/319 [=====] - 3s 10ms/step - loss: 0.4560 - recall: 0.7773 - val_loss: 0.4325 - val_recall: 0.6840

Epoch 11/50
319/319 [=====] - 3s 10ms/step - loss: 0.4552 - recall: 0.7771 - val_loss: 0.4415 - val_recall: 0.7025

Epoch 12/50
319/319 [=====] - 3s 11ms/step - loss: 0.4483 - recall: 0.7867 - val_loss: 0.4410 - val_recall: 0.6871

Epoch 13/50
319/319 [=====] - 3s 11ms/step - loss: 0.4448 - recall: 0.7786 - val_loss: 0.4358 - val_recall: 0.6810

Epoch 14/50
319/319 [=====] - 4s 12ms/step - loss: 0.4447 - recall: 0.7824 - val_loss: 0.4468 - val_recall: 0.7025

Epoch 15/50
319/319 [=====] - 3s 11ms/step - loss: 0.4397 - recall: 0.7741 - val_loss: 0.4504 - val_recall: 0.7025

Epoch 16/50
319/319 [=====] - 4s 13ms/step - loss: 0.4415 - recall: 0.7926 - val_loss: 0.4386 - val_recall: 0.6871

Epoch 17/50
319/319 [=====] - 3s 10ms/step - loss: 0.4402 - recall: 0.7910 - val_loss: 0.4432 - val_recall: 0.6994

Epoch 18/50
319/319 [=====] - 6s 20ms/step - loss: 0.4388 - recall: 0.7985 - val_loss: 0.4200 - val_recall: 0.6656

Epoch 19/50
319/319 [=====] - 8s 24ms/step - loss: 0.4336 - recall: 0.7922 - val_loss: 0.4345 - val_recall: 0.6994

Epoch 20/50
319/319 [=====] - 8s 24ms/step - loss: 0.4347 - recall: 0.7938 - val_loss: 0.4370 - val_recall: 0.6902

Epoch 21/50
319/319 [=====] - 8s 25ms/step - loss: 0.4339 - recall: 0.79
67 - val_loss: 0.4335 - val_recall: 0.6963

Epoch 22/50
319/319 [=====] - 8s 26ms/step - loss: 0.4285 - recall: 0.79
42 - val_loss: 0.4196 - val_recall: 0.6871

Epoch 23/50
319/319 [=====] - 8s 25ms/step - loss: 0.4273 - recall: 0.80
00 - val_loss: 0.4333 - val_recall: 0.7025

Epoch 24/50
319/319 [=====] - 8s 25ms/step - loss: 0.4280 - recall: 0.80
53 - val_loss: 0.4215 - val_recall: 0.6902

Epoch 25/50
319/319 [=====] - 5s 16ms/step - loss: 0.4291 - recall: 0.80
47 - val_loss: 0.4453 - val_recall: 0.7147

Epoch 26/50
319/319 [=====] - 8s 25ms/step - loss: 0.4278 - recall: 0.81
16 - val_loss: 0.4298 - val_recall: 0.7025

Epoch 27/50
319/319 [=====] - 9s 27ms/step - loss: 0.4285 - recall: 0.80
57 - val_loss: 0.4273 - val_recall: 0.6963

Epoch 28/50
319/319 [=====] - 6s 20ms/step - loss: 0.4243 - recall: 0.80
93 - val_loss: 0.4317 - val_recall: 0.7117

Epoch 29/50
319/319 [=====] - 7s 23ms/step - loss: 0.4252 - recall: 0.80
87 - val_loss: 0.4392 - val_recall: 0.7025

Epoch 30/50
319/319 [=====] - 7s 23ms/step - loss: 0.4246 - recall: 0.80
24 - val_loss: 0.4364 - val_recall: 0.7178

Epoch 31/50
319/319 [=====] - 7s 22ms/step - loss: 0.4257 - recall: 0.80
71 - val_loss: 0.4389 - val_recall: 0.7117

Epoch 32/50
319/319 [=====] - 3s 10ms/step - loss: 0.4250 - recall: 0.81
30 - val_loss: 0.4315 - val_recall: 0.7025

Epoch 33/50
319/319 [=====] - 5s 15ms/step - loss: 0.4225 - recall: 0.79
85 - val_loss: 0.4384 - val_recall: 0.7086

Epoch 34/50
319/319 [=====] - 6s 19ms/step - loss: 0.4209 - recall: 0.80
65 - val_loss: 0.4274 - val_recall: 0.6779

Epoch 35/50
319/319 [=====] - 8s 25ms/step - loss: 0.4197 - recall: 0.80
59 - val_loss: 0.4246 - val_recall: 0.6902

Epoch 36/50
319/319 [=====] - 8s 25ms/step - loss: 0.4164 - recall: 0.81
81 - val_loss: 0.4248 - val_recall: 0.6810

Epoch 37/50
319/319 [=====] - 7s 22ms/step - loss: 0.4183 - recall: 0.80
73 - val_loss: 0.4263 - val_recall: 0.6871

Epoch 38/50
319/319 [=====] - 9s 27ms/step - loss: 0.4139 - recall: 0.81
14 - val_loss: 0.4281 - val_recall: 0.6810

Epoch 39/50
319/319 [=====] - 7s 23ms/step - loss: 0.4192 - recall: 0.81
28 - val_loss: 0.4310 - val_recall: 0.7117

Epoch 40/50
319/319 [=====] - 7s 21ms/step - loss: 0.4202 - recall: 0.81
51 - val_loss: 0.4360 - val_recall: 0.6871

```

Epoch 41/50
319/319 [=====] - 8s 25ms/step - loss: 0.4133 - recall: 0.81
28 - val_loss: 0.4284 - val_recall: 0.6902
Epoch 42/50
319/319 [=====] - 6s 19ms/step - loss: 0.4159 - recall: 0.81
42 - val_loss: 0.4288 - val_recall: 0.6748
Epoch 43/50
319/319 [=====] - 7s 23ms/step - loss: 0.4126 - recall: 0.80
79 - val_loss: 0.4345 - val_recall: 0.6933
Epoch 44/50
319/319 [=====] - 8s 24ms/step - loss: 0.4166 - recall: 0.81
50 - val_loss: 0.4242 - val_recall: 0.6748
Epoch 45/50
319/319 [=====] - 5s 17ms/step - loss: 0.4159 - recall: 0.81
57 - val_loss: 0.4263 - val_recall: 0.6810
Epoch 46/50
319/319 [=====] - 8s 24ms/step - loss: 0.4128 - recall: 0.81
14 - val_loss: 0.4270 - val_recall: 0.6779
Epoch 47/50
319/319 [=====] - 7s 22ms/step - loss: 0.4133 - recall: 0.81
55 - val_loss: 0.4240 - val_recall: 0.6994
Epoch 48/50
319/319 [=====] - 7s 23ms/step - loss: 0.4199 - recall: 0.81
16 - val_loss: 0.4270 - val_recall: 0.7055
Epoch 49/50
319/319 [=====] - 8s 24ms/step - loss: 0.4148 - recall: 0.81
42 - val_loss: 0.4111 - val_recall: 0.6748
Epoch 50/50
319/319 [=====] - 8s 26ms/step - loss: 0.4120 - recall: 0.81
71 - val_loss: 0.4316 - val_recall: 0.7025


```

Loss function

```

In [ ]: #Plotting Train Loss vs Validation Loss
plt.plot(history_5.history['loss'])
plt.plot(history_5.history['val_loss'])
plt.title('model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

```



```
In [ ]: #Plotting Train recall vs Validation recall
plt.plot(history_5.history['recall'])
plt.plot(history_5.history['val_recall'])
plt.title('model recall')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```




```
In [ ]: y_train_pred = model_5.predict(X_train_smote)
        #Predicting the results using 0.5 as the threshold
        y_train_pred = (y_train_pred > 0.5)
        y_train_pred
```

319/319 [=====] - 6s 14ms/step

```
Out[ ]: array([[ True],
               [False],
               [False],
               ...,
               [ True],
               [ True],
               [ True]])
```

```
In [ ]: y_val_pred = model_5.predict(X_val)
        #Predicting the results using 0.5 as the threshold
        y_val_pred = (y_val_pred > 0.5)
        y_val_pred
```

50/50 [=====] - 2s 16ms/step

```
Out[ ]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [ True],
               [ True]])
```

```
In [ ]: model_name = "NN with SMOTE,Adam & Dropout"
```

```
train_metric_df.loc[model_name] = recall_score(y_train_smote,y_train_pred)
valid_metric_df.loc[model_name] = recall_score(y_val,y_val_pred)
```

Classification report

```
In [ ]: cr=classification_report(y_train_smote,y_train_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.83	0.84	0.83	5096
1	0.84	0.82	0.83	5096
accuracy			0.83	10192
macro avg	0.83	0.83	0.83	10192
weighted avg	0.83	0.83	0.83	10192

```
In [ ]: #classification report
cr=classification_report(y_val,y_val_pred)
print(cr)
```


	precision	recall	f1-score	support
0	0.92	0.83	0.87	1274
1	0.51	0.70	0.59	326
accuracy			0.80	1600
macro avg	0.72	0.77	0.73	1600
weighted avg	0.83	0.80	0.81	1600

Confusion matrix

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_train_smote, y_train_pred)
```



```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_val,y_val_pred)
```



Model Performance Comparison and Final Model Selection

```
In [ ]: print("Training performance comparison")
train_metric_df
```

Training performance comparison

```
Out[ ]:
```

	recall
NN with SGD	0.000767
NN with Adam	0.523006
NN with Adam & Dropout	0.603528
NN with SMOTE & SGD	0.766287
NN with SMOTE & Adam	0.940542
NN with SMOTE,Adam & Dropout	0.822998

```
In [ ]: print("Validation set performance comparison")
        valid_metric_df
```

Validation set performance comparison

```
Out[ ]:
```

	recall
NN with SGD	0.003067
NN with Adam	0.469325
NN with Adam & Dropout	0.490798
NN with SMOTE & SGD	0.714724
NN with SMOTE & Adam	0.604294
NN with SMOTE,Adam & Dropout	0.702454

```
In [ ]: train_metric_df - valid_metric_df
```

```
Out[ ]:
```

	recall
NN with SGD	-0.002301
NN with Adam	0.053681
NN with Adam & Dropout	0.112730
NN with SMOTE & SGD	0.051563
NN with SMOTE & Adam	0.336247
NN with SMOTE,Adam & Dropout	0.120544


```
In [ ]: y_test_pred = model_3.predict(X_test)
        y_test_pred = (y_test_pred > 0.5)
        print(y_test_pred)
```

63/63 [=====] - 0s 4ms/step
[[False]
[True]
[False]
...
[True]
[False]
[False]]

```
In [ ]: #Lets print classification report
cr=classification_report(y_test,y_test_pred)
print(cr)
```

	precision	recall	f1-score	support
0	0.92	0.71	0.80	1593
1	0.40	0.75	0.53	407
accuracy			0.72	2000
macro avg	0.66	0.73	0.66	2000
weighted avg	0.81	0.72	0.75	2000

```
In [ ]: #Calculating the confusion matrix
make_confusion_matrix(y_test,y_test_pred)
```



Actionable Insights and Business Recommendations

Recommendations and Insights

- The EDA showed that most predictors were not strong and were unbalanced
- The models that used SMOTE performed the best with the highest recall values. SMOTE balanced the data set and took care of underrepresented classes. The imbalanced data led to weaker models for the first 3 tested models. ### Conclusion
- The Neural Network Model with SGD and SMOTE is the model with the highest recall. (Model 3)
- This model was slightly overfit but produced the highest recall at 71.5% in the validation set with slight overfitting.

- This model is the best to predict customers action since we are maximizing recall and minimizing false negatives.

By Raghuram Palaniappan